

Gathering Bibliometric Information from the Scopus API using the `rscopus` R Package

by John Muschelli

Abstract We demonstrate how to download author and affiliation using the R package `rscopus`, interacting with the Elsevier Scopus API. We demonstrate how to manipulate the output from the API into organized data that can be analyzed. We present options on how to obtain the number of citations from an author and use that to calculate citation metrics with the output. We also show how to download additional information from an article such as the figures and supplemental material.

Introduction

Elsevier provides a number of APIs (application program interfaces) to extract information about scholarly publications (<https://dev.elsevier.com/>). This information can provide insights into how research is done and provide the data to perform experiments and analyses of metascience. Elsevier's Scopus a repository of information about articles, authors, and institutions. Scopus claims to have the largest database of this information (<https://www.elsevier.com/solutions/scopus>). Therefore, providing users an interface to this repository should be worthwhile.

One common task for researchers is to keep his or her curriculum vitae (CV) up to date. That requires having accurate information on the papers published and under submission. Keeping track of these papers can be tedious and solutions could exist if one could retrieve papers automatically. One concern is missing certain crucial papers in your CV. Although `rscopus` does not provide these tools specifically, it provides tools to download publications about an author, which can be used to consistently cross-reference information about publications with a CV.

Additionally on CVs, one may present information of the impact of a paper. This can be done by highlighting certain pieces of information, such as done in NIH biosketches, or ranking papers based on some metric. As promotions and grant review may be affected by highlighting impact, these metrics can have real-life implications. One metric commonly used is the number of citations. Additionally, information about the journal impact factor may be shown. We do not imply that these are particularly good metrics or metrics that reflect true impact, but are simply those that we have seen used in practice.

The `rscopus` package allows users to interface with Scopus APIs and gather information about authors, affiliations, articles, and abstracts. Currently in R, packages exist for bibliometric analysis, but commonly require the data to be downloaded from a website or online interface. For example, the `bibliometrix` (2) package provides a level of integration that is useful for using multiple packages that deal with bibliometric data, incorporating functionality from `rscopus`. The `bibliometrix` package also enables users to analyze data from ISI Web of Knowledge (WoK) and PubMed.

Web of Knowledge is one competitor to Scopus, but `bibliometrix` does not have an interface to the WoK API; therefore data must be manually exported from the site into R. Additional access to the web of Science API would be useful and has been implemented in a GitHub package `rwos` (<https://github.com/juba/rwos>), but is not on CRAN and has not been updated recently (over 1 year). We have also created an interface with the Web of Science APIs (<https://github.com/muschellij2/webofscience>), but have not been given access to the APIs to test them as it is still in beta.

As compared to Google Scholar, Scopus and WoK claim the information from these sources is more curated. Other packages such as `scholar` (7) and `gcite` (8) can provide interfaces to the Google Scholar citation information. Using these in combination with `rscopus` can more likely guarantee complete information. Also, comparing these different repositories can demonstrate biases or differences in the number of citations across the different platforms. Also, Scopus has a true API compared to the packages above that extract information from the Google Scholar webpage, which will likely be more stable, more robust to API changes, and be within the terms of service for the user.

Scopus has a number of APIs available (https://dev.elsevier.com/sc_apis.html). Here we will present examples of how to use the `rscopus` package, including: searching for authors or affiliations, calculating citation indices for an author, retrieving abstract information on articles, and downloading artifacts from an article.

API Key

Before using the package, one must obtain an access key to the API from [Elsevier](#) with the following steps:

1. Go to <https://dev.elsevier.com/user/login>. Login or create a free account.
2. Click "Create API Key". Put in a label, such as "rscopus key". Add a website. <http://example.com> is fine if you do not have a site.
3. **Read** and agree to the terms of service if you do indeed agree.
4. You should be able to test out the API key using the [interactive Scopus APIs](#).
5. Add `Elsevier_API = "API KEY GOES HERE"` to `~/.Renvi`ron file, or add `export Elsevier_API=API KEY GOES HERE` to your `~/.bash_profile`.

Alternatively, you can either set the API key using `rscopus::set_api_key`, which will implicitly set `options("elsevier_api_key" = api_key)`. You can access the API key using `rscopus::get_api_key`.

Once you have an API key set up, you can access the key using `get_api_key`, which will check multiple places for the presence of the key:

```
library(rscopus)
key = get_api_key()
```

As you may want to hide this key from documents, the default print options do not display it:

```
key
```

```
<hidden api key, use print(, reveal = TRUE) to see it>
```

There are other helpful functions such as `have_api_key`, which tests if an API key has been specified:

```
have_api_key()
```

```
[1] TRUE
```

In each function to an API endpoint, you may pass in the API key as the argument `api_key` as well, rather than specifically look in the `~/.Renvi`ron or `options()`.

A note about API keys and IP addresses

The API Key is bound to a set of IP addresses, usually from your institution or organization (see https://dev.elsevier.com/tecdoc_api_authentication.html). Therefore, if you are using **rscopus** for a Shiny application, you must host the Shiny application from the institution/organization servers in some way. Also, you cannot access the Scopus API with this key if you are offsite and must VPN into the server or use a computing cluster with an institution IP.

Methods: Use cases

In this section we will present some use cases for the API, focusing on retrieving information about authors and affiliations.

Processing author names to identifiers

Researchers commonly would like to gather information about a set of authors. Most times the authors are given by first and last names or initials; additional information such as affiliation may be available. Scopus provides unique identifier for authors (`au_id`) or affiliations (`affil_id`), amongst others. In many cases with the API, you will specify the author identifier (`au_id`) instead of a first and last name, as there may be many authors with the same name. In order to get the author identifier from Scopus, you can search using a first and last name using the `process_author_name` command. For example, let us try to identify the author ID for John Muschelli:

```
auth_info = process_author_name(last_name = "Muschelli", first_name = "John",
                               verbose = FALSE)
auth_info
```

```
$au_id
[1] "40462056100"
```

```
$first_name
[1] "John"
```

```
$last_name
[1] "Muschelli"
```

The output is a simple list of first and last name with an author ID. The function chooses the first author found, which may be useful if the author name is somewhat unique. We will show below how to search when the name is not as unique. This ID is unique to this author, though curation errors do happen and someone may have 2 unique identifiers. These identifiers can be [merged by request](#) on the Scopus website.

Retrieving author citation data

In order to get data about papers and citations from an author, the `author_data` function will retrieve this information:

```
jm = author_data(last_name = "Muschelli", first_name = "John", verbose = FALSE)
```

We see the output is a list of the converted entries from the JSON output Scopus API, a `data.frame` of the results for citations, and a list named `full_data`:

```
class(jm)
[1] "list"

names(jm)
[1] "entries"      "df"           "first_name"   "last_name"    "au_id"
[6] "full_data"
```

The `data.frame` `df` has the information many users wish to retrieve, which is information about the author documents such as the Scopus ID for that paper (`dc:identifier`), the title of the paper (`dc:title`), and the number of citations (`citedby-count`). Here we present `short_title`, first 3 “relevant” words of the title (see `unique_title` in supplemental material), instead of the full document title from `dc:title` for viewing purposes as titles can be quite long.

```
jm$df$short_title = unique_title(jm$df$`dc:title`)
head(jm$df[, c("dc:identifier", "short_title", "citedby-count")])
```

	dc:identifier	short_title	citedby-count
1	SCOPUS_ID:85053246791	Objective Evaluation Multiple	1
2	SCOPUS_ID:85043338865	MIMOSA: Automated Method	3
3	SCOPUS_ID:85047750078	Radiomic subtyping improves	8
4	SCOPUS_ID:85028874240	Feasibility Coping Effectiveness	2
5	SCOPUS_ID:85050271095	Freesurfer: Connecting Freesurfer	0
6	SCOPUS_ID:85055916029	dual modeling approach	0

We see that the `full_data` has the object `df` duplicated in it, as well as other `data.frames`:

```
names(jm$full_data)
[1] "df"           "affiliation"  "author"
```

These `data.frames` can have additional information about co-author affiliations or co-author information. This information may be useful for creating network graphs. For example, to get all authors from all the papers, you can use the author element from `full_data`:

```
head(jm$full_data$author[, c("authid", "authname", "surname", "afid.$", "entry_number")])
```

	authid	authname	surname	afid.\$	entry_number
1	8431704700	Commowick O.	Commowick	60030553	1
2	57203861434	Istace A.	Istace	60001780	1
3	57199507814	Kain M.	Kain	60030553	1
4	57197801981	Laurent B.	Laurent	60105610	1
5	57203867656	Leray F.	Leray	60030553	1
6	57203864793	Simon M.	Simon	60030553	1

The column `entry_number` indicates which element this information came in the entry list (retrieved from content from `httr` (11), which calls fromJSON from `jsonlite` (9)). This column should merge with the `data.frame` of citations, as well as the information about author affiliations, which is located in the `affiliation.data.frame` from `full_data`:

```
head(jm$full_data$affiliation[, c("afid", "affiliation-country", "entry_number", "affilname")])
```

	afid	affiliation-country	entry_number		affilname
1	60030553	France	1		Universite de Rennes 1
2	60001780	France	1		Centre Hospitalier Lyon-Sud
3	60105610	France	1		Laboratoire de Traitement de l'Information Médicale
4	60062760	France	1		Centre de Recherche en Acquisition et Traitement d'Images pour la Santé
5	60028893	France	1		Centre Hospitalier Universitaire de Rennes
6	60028893	France	1		Centre Hospitalier Universitaire de Rennes

This information is rich for understanding information about an author's publication record, how many citations are recorded for a specific article, which journals have been published in, and who has co-authored publications with an author. Here we rank the authors based on number of co-authored publications. We will use `dplyr` (13) for data manipulation:

```
library(dplyr)
au_id = unique(jm$df$au_id)
co_authors = jm$full_data$author %>%
  filter(!authid %in% au_id)
co_authors = co_authors %>%
  add_count(authid) %>%
  select(n, authid, authname, surname) %>%
  distinct() %>% arrange(-n)
head(co_authors)

# A tibble: 6 x 4
   n authid      authname      surname
<int> <chr>      <chr>      <chr>
1    18 7003584880 Mostofsky S. Mostofsky
2    16 7003771238 Pekar J.    Pekar
3    12 7103268788 Hanley D.   Hanley
4    11 55587557800 Sweeney E.  Sweeney
5    10 8582425200 Barber A.   Barber
6     9 6603085071 Crainiceanu C. Crainiceanu
```

More analyses of co-authorship can be done, such as determining how many international collaborators an author had on papers based on affiliation.

Retrieving summary information about an author

The `author_retrieval` function can gather summary information about an author using the author identifier or name.

```
author_info = author_retrieval(last_name = "Muschelli", first_name = "J")

      auth_name      au_id affil_id
1 John Muschelli 40462056100 60006183
      affil_name
1 Johns Hopkins Bloomberg School of Public Health

names(author_info$content)

[1] "author-retrieval-response"
```

```
class(author_info$content$`author-retrieval-response`)
[1] "list"
```

In the standard output from the Scopus API after conversion in **httr**, there are elements of the list named `entries` or `entry`. The low-level function `gen_entries_to_df` attempts to coerce this list into a standard `data.frame` for more usability, but may not perform perfectly as lists from JSON cannot always be directly coerced into a rectangular format. For example, here we will convert that output into a `data.frame`:

```
gen_entries_to_df(author_info$content$`author-retrieval-response`)$df
  @status @_fa
1   found true
                                coredata.prism:url
1 http://api.elsevier.com/content/author/author_id/40462056100
  coredata.dc.identifier   coredata.eid   coredata.orcid
1  AUTHOR_ID:40462056100 9-s2.0-40462056100 0000-0001-6469-1750
  coredata.document-count coredata.cited-by-count coredata.citation-count
1                        37                        766                        914
                                coredata.link.@href
1 https://www.scopus.com/authid/detail.uri?partnerID=Hz0xMe3b&authorId=40462056100&origin=inward
  coredata.link.@rel coredata.link.@_fa preferred-name.surname
1   scopus-author      true      Muschelli
  preferred-name.given-name preferred-name.initials
1             John             J.
  affiliation-current.affiliation-name
1 Johns Hopkins Bloomberg School of Public Health
  affiliation-current.affiliation-city
1             Baltimore
  affiliation-current.affiliation-country publication-range.start
1             United States             2011
  publication-range.end entry_number
1             2018             1
```

but this list from `author_retrieval` typically only has one element, and may be easily referenced using `$` as a list. Most of the functions implicitly call `gen_entries_to_df` to convert the data when available, and additional organizing functions exist such as `process_complete_multi_author_info` discussed below. Overall, `author_retrieval` gives general info about an author, but this information and more can likely be extracted using the `author_data` function above.

Calculating author indices

With the data from the `author_data` output (`jm$df`), we can calculate citation indices. Here we calculate the overall *h*-index (6):

```
h_data = jm$df %>%
  mutate(citations = as.numeric(`citedby-count`)) %>%
  arrange(-citations) %>%
  mutate(n_papers = 1:n())
head(h_data[, c("short_title", "citations", "n_papers")])
      short_title citations n_papers
1   Minimally invasive surgery    151      1
2 Reduction motion-related artifacts    74      2
3 Disruption functional organization    71      3
4 Large-scale radiomic profiling    60      4
5 Minimally invasive evacuation    60      5
6 Thrombolytic removal intraventricular    59      6

h_index = max(which(h_data$citations >= h_data$n_papers))
cat(paste0("Calculated h-index is ", h_index))

Calculated h-index is 16
```

Using **ggplot2** (10), we can also visually show the *h*-index computation, where we plot the number of citations versus the number of papers (cumulatively) along with the X-Y line:




Figure 1: Calculating an h-index. Here we plot the number of papers versus the number of citations for that paper. This plot is the basis for the h-index. The X=Y line is displayed in black and the red line is where the curve passes the X=Y line, which is the h-index, a value of 16.

```
library(ggplot2)
h_data %>%
  ggplot(aes(x = n_papers, y = citations)) +
    geom_point() + geom_abline(slope = 1, intercept = 0) +
    geom_hline(yintercept = h_index, color = "red")
```

Additional indices can be created from the data, such as the g-index (4):

```
h_data = h_data %>% mutate(sum_citations = cumsum(citations))
g_index = max(which(h_data$sum_citations >= h_data$n_papers^2))
cat(paste0("Calculated g-index is ", g_index))
```

Calculated g-index is 30

We could create a series of metrics based on this data and compare it to or supplement other metrics of performance. Note, however, longitudinal information is lacking on the number of citations per year. We will discuss citation information over time in Section 2.2.7.

Retrieving information about an author

In `process_author_name`, we demonstrated how to get information from an author with a relatively unique name. If this is not the case, the `get_complete_author_info`, which is the backend function for `process_author_name`, can present more results. In order to retrieve author IDs from first and last names, the `get_complete_author_info` can be used. Here we search for authors with the last name West and first initial M:

```
last_name = "West"
first_name = "M"
auth_info_list = get_complete_author_info(last_name = last_name, first_name = first_name)
class(auth_info_list)

[1] "list"

names(auth_info_list)
```

```
[1] "get_statement" "content"
```

We see here, which is common in some low-level functions returned from the rscopus API, the output is a list with elements `get_statement`, which returns an object of class `response` (from the **http** package), and `content`, which is the content from the response. Most times, the content is of interest, but failed requests may be explored with the `get_statement` output for debugging.

Again, we can use `gen_entries_to_df` convert that output into a `data.frame`, which is more manageable. Here show see the author information retrieved from the call:

```
coerced = gen_entries_to_df(auth_info_list$content$search-results$entry)
names(coerced)

[1] "df"          "name-variant" "subject-area"

head(coerced$df[, c("dc:identifier", "preferred-name.surname",
                    "preferred-name.given-name", "affiliation-current.affiliation-name" )])

      dc:identifier preferred-name.surname preferred-name.given-name
1 AUTHOR_ID:35480328200                West Catharine
2 AUTHOR_ID:35419377800                West Malcolm J.
3 AUTHOR_ID:7003392768                 Diener-West Marie
4 AUTHOR_ID:7402395730                West Robert M.
5 AUTHOR_ID:7402068812                West Michael Abigail
6 AUTHOR_ID:7401998578                West David M.

      affiliation-current.affiliation-name
1 University of Manchester
2 James Cook University, Australia
3 Johns Hopkins Bloomberg School of Public Health
4 University of Leeds
5 University of Pittsburgh Medical Center
6 Massey University
```

We see this has information about the multiple authors returned, along with names, variations on those names, number of documents, and affiliations. We can then extract the author ID we want from this `data.frame`. This process is wrapped in the `get_author_info`:

```
auth_info_df = get_author_info(last_name = last_name,
                              first_name = first_name)
head(auth_info_df)

      auth_name      au_id affil_id
1 Catharine West 35480328200 60003771
2 Malcolm J. West 35419377800 60019870
3 Marie Diener-West 7003392768 60006183
4 Robert M. West 7402395730 60012070
5 Michael Abigail West 7402068812 60012018
6 David M. West 7401998578 60008221

      affil_name
1 University of Manchester
2 James Cook University, Australia
3 Johns Hopkins Bloomberg School of Public Health
4 University of Leeds
5 University of Pittsburgh Medical Center
6 Massey University
```

We should note this information is condensed and a subset that is available from `get_complete_author_info`, also with more standardized column names. Now, one could use the `au_id` retrieved from this output for searching and retrieval.

If you had an affiliation ID, such as 60006183 for the Johns Hopkins Bloomberg School of Public Health, we can pass this to `get_author_info` or `process_author_name`:

```
spec_affil = get_author_info(
  last_name = last_name,
  first_name = first_name,
  affil_id = 60006183)
spec_affil
```

```

      auth_name      au_id affil_id
1 Marie Diener-West 7003392768 60006183
      affil_name
1 Johns Hopkins Bloomberg School of Public Health

```

Thus, we can use the combination of the names of the author and affiliation information. This may be useful when searching multiple authors at the same institution. We could also pass in `affil_name` instead of `affil_id` to `get_author_info`, but the name must be extremely specific for institution as the function will call `get_affiliation_info`, which is discussed below.

Retrieving information about multiple authors

In order to get information from multiple authors, one could loop over author information, but this is inefficient for code and API calls. The `complete_multi_author_info` function can perform this operation. One caveat is that it requires author identifiers and not names. We can take the author IDs from `auth_info_df` to retrieve information for all these authors:

```

all_author_info = complete_multi_author_info(au_id = auth_info_df$au_id)
names(all_author_info)

```

```
[1] "get_statement" "content"      "au_id"
```

This result is again a low-level output from the API. We can use the `process_complete_multi_author_info` function to process this into a more amenable solution:

```

processed = process_complete_multi_author_info(all_author_info)
head(names(processed))

[1] "35480328200" "35419377800" "7003392768"  "7402395730"  "7402068812"
[6] "7401998578"

```

Now, each element is the author ID, which contains a list of data.frames.

For example, for the ID 35480328200, we can see the elements for that specific author.

```

names(processed$`35480328200`)

[1] "info"          "links"          "other_names"
[4] "subject_areas" "affiliation_history" "journals"

```

We can then access the affiliation history, which includes information such as department and university:

```

head(processed$`35480328200`$affiliation_history[, c("ip-doc.afdispname", "ip-doc.id", "ip-doc.type"), 3])

      ip-doc.afdispname ip-doc.id
1 University of Manchester, Division of Cancer Sciences 119995430
2 University of Manchester 60003771
3 University of Manchester, Radiotherapy Related Research 110273198
ip-doc.type
1 dept
2 parent
3 dept

```

The `multi_author_info` will perform both of these operations together. This result is still not “tidy” (14) in many respects, but parts can be combined using `purrr` (5):

```

journals = purrr::map_df(processed, `$$`, "journals", .id = "au_id")
head(journals)

      au_id type      sourcetitle
1 35480328200 j      Cancer Letters
2 35480328200 j Journal of Cancer Research and Clinical Oncology
3 35480328200 j      Nature Reviews Cancer
4 35480328200 j      Scientific Reports
5 35480328200 j Annals of the Royal College of Surgeons of England
6 35480328200 j      Cancer Letters
      sourcetitle-abbrev issn

```



```

1          Cancer Lett. 03043835
2 J. Cancer Res. Clin. Oncol. 01715216
3          Nat. Rev. Cancer 1474175X
4          Sci. Rep. 20452322
5 Ann. R. Coll. Surg. Engl. 00358843
6          Cancer Lett. 18727980

```

Thus, in the above code we could compare the journals each author published in compared to the other authors in the list.

Citations over time

Some APIs from Elsevier are disabled by default (see https://dev.elsevier.com/api_key_settings.html). Notably, the Citations Overview API is disabled, which allows users to access information about citations over time for articles of authors. This citation information is particularly useful for creating bibliometric indices that rely on subsetting data based on minimum or maximum year or calculating time-dependent metrics. The `rscopus` package interfaces with these APIs, but the API must be enabled for that specific API key. Here, we show the output of the failed request as it may be useful for debugging:

```

sc_id = jm$df$`dc:identifier`[1]
cit_try = citation_retrieval(scopus_id = sc_id)
httr::status_code(cit_try$get_statement)

[1] 403

cit_try$content

$`service-error`
$`service-error`$status
$`service-error`$status$statusCode
[1] "AUTHENTICATION_ERROR"

$`service-error`$status$statusText
[1] "Requestor configuration settings insufficient for access to this resource."

```

On the Scopus website, however, one can searching for authors, select up to 15 authors, and then create a “Citation Overview”, which will give this citation information, and export it in a CSV format. The `rscopus` package provides a `read_cto` function to read in this data.

We also provide an example export from a single author in the package:

```

file = system.file("extdata", "CTOExport.csv", package = "rscopus")
citations_over_time = rscopus::read_cto(file)
names(citations_over_time)

[1] "data"          "year_columns"  "author_information"

```

The real information is in the data element of this list, which contains the yearly citation information (again titles have been shortened):

```

yr_cols = citations_over_time$year_columns
citations_over_time = citations_over_time$data
citations_over_time = citations_over_time %>%
  mutate(short_title = unique_title(`Document Title`))
head(citations_over_time[, c("short_title", yr_cols[1:5])])

```

	short_title	<2008	2008	2009	2010	2011
1	Objective Evaluation Multiple	0	0	0	0	0
2	MIMOSA: Automated Method	0	0	0	0	0
3	Radiomic subtyping improves	0	0	0	0	0
4	Feasibility Coping Effectiveness	0	0	0	0	0
5	Freesurfer: Connecting Freesurfer	0	0	0	0	0
6	Thrombolytic removal intraventricular	0	0	0	0	0

In the citation overview, you must specify a range of years on Scopus, with a maximum of 15 years. As many times this wide format is not what you want to plot, or in a “tidy” format, the `rscopus` helper function `read_cto_long` will read the data in long format, done by `tidyr` (12). Here we use `dplyr` to arrange the data by maximum number of citations per year:

```

long_cite = rscopus::read_cto_long(file)
long_cite = long_cite$data %>%
  group_by(`Document Title`, year) %>% # get the citations per year
  summarize(citations = sum(citations), # aggregate - some duplicates merged
    `Publication Year` = unique(`Publication Year`)) %>% # keep the year in data
  mutate(short_title = unique_title(`Document Title`))
long_cite = long_cite %>% arrange(-citations, year, short_title)
head(long_cite[, c("short_title", "year", "citations")])

```

```

# A tibble: 6 x 3
  short_title          year citations
  <chr>              <fct>    <int>
1 Minimally invasive surgery 2015      36
2 Minimally invasive surgery 2017      35
3 ISLES 2015 public          2018      28
4 Large-scale radiomic profiling 2018      26
5 Thrombolytic removal intraventricular 2018      25
6 Minimally invasive surgery 2016      24


```

Thus, we have one record per year and article. Here we will plot the cumulative citations per each paper over the years of publication and label the top 3 cited papers:

```

# get cumulative sum
csum = long_cite %>%
  # any missing data had no citations
  mutate(citations = ifelse(is.na(citations), 0, citations)) %>%
  arrange(`Document Title`, year) %>% # sort for cumsum
  group_by(`Document Title`) %>%
  mutate(citations = cumsum(citations))
# remove past and future with as.integer
csum = csum %>%
  mutate(year = as.integer(as.character(year))) %>%
  filter(!is.na(year)) %>% # remove < 2008 and > 2018 years
  filter(year >= `Publication Year` & citations > 0) # keep only relevant data for paper
# grab last citations and top 3 papers
last_year = csum %>%
  arrange(`Document Title`, year) %>% # sort for slice later
  group_by(`Document Title`) %>%
  slice(n()) %>% # keep last as max citations
  ungroup %>% arrange(-citations) %>%
  head(3) # get top 3
g = ggplot(csum,
  aes(x = year, y = citations, color = short_title )) +
  xlim(c(2010, 2018)) + geom_line() +
  # label the titles numbers for top 3
  geom_text(data = last_year, size = 3, aes(label = short_title),
    nudge_x = -1, nudge_y = 5)
# don't want label for document title - too many entries
g + guides(color = FALSE) + theme(text = element_text(size = 20))

```



Thus, we can present visually how the number of citations has changed and may be able to highlight which papers are gaining or waning in citations over time. This plot may not provide deep insights, but the same plot could be made by scientific sub-fields or for specific journals, which may indicate trends in published articles for example.

Retrieving affiliation information

In order to get information about an affiliation, the `get_affiliation_info` can be used. Here we will look for the pattern Johns Hopkins:

```
jhu_info = get_affiliation_info(affil_name = "Johns Hopkins")
head(jhu_info[, c("affil_id", "affil_name")])
```

	affil_id	affil_name
1	60005248	Johns Hopkins University
2	60001117	The Johns Hopkins School of Medicine
3	60006183	Johns Hopkins Bloomberg School of Public Health
4	60001555	Johns Hopkins Hospital
5	60003443	Johns Hopkins Medical Institutions
6	60022054	The Johns Hopkins University Applied Physics Laboratory

This function implicitly calls `affil_search`, a lower-level function which searches the affiliation information from Scopus. Additional information can be extracted using `affil_search`, but this typically includes a large number of records as it searches all the documents. This affiliation ID again can be used to be more specific when searching authors or documents. Thus, we see that if we had only used "Johns Hopkins" as the query and not "Johns Hopkins Bloomberg School of Public Health", we would get back a different top candidate. This specificity is why we recommend using affiliation ID if using affiliation information in `get_author_info`.

Searching articles by abstract

In some cases, one may have an article in mind and would like information about the authors of that article. In order to get the author IDs from the article identifier, one can use the `abstract_retrieval`

function. There are multiple identifiers that can be used, such as PubMed ID, DOI, and we will use the Scopus ID as that is what is returned from the `author_data` output:

```
sc_id = jm$df$`dc:identifier`[1]
# retrieve abstract
abstract = abstract_retrieval(id = sc_id, identifier = "scopus_id")
```

Here we will extract the abstract information from the result:

```
sc_info = abstract$content$`abstracts-retrieval-response`
substr(sc_info$coredata$`dc:description`, 1, 76)
```

```
[1] "© 2018, The Author(s). We present a study of multiple sclerosis segmentation"
```

Here we can extract information about the authors of the paper:

```
sc_df = purrr::map_df(sc_info$authors[[1]],
  as.data.frame, stringsAsFactors = FALSE, make.names = FALSE)
head(sc_df[, c("ce.given.name", "ce.initials", "X.auid")])
```

	ce.given.name	ce.initials	X.auid
1	Olivier	O.	8431704700
2	Audrey	A.	57203861434
3	Michaël	M.	57199507814
4	Baptiste	B.	57197801981
5	Florent	F.	57203867656
6	Mathieu	M.	57203864793

This information is located within the `author_data.frame` from the `full_data` as well. Note, however that the information from `author_data` was obtained by searching Scopus for an author, whereas the information from `abstract_retrieval` was obtained by searching for a specific paper. As we took the first entry from the Scopus identifier, we will subset the author data by `entry_number` 1 from the `author_data.frame` to show the relevant info matches:

```
paper_author_info = jm$full_data$author
head(paper_author_info[paper_author_info$entry_number == 1, c("authid", "authname", "surname")])
```

	authid	authname	surname
1	8431704700	Commowick O.	Commowick
2	57203861434	Istace A.	Istace
3	57199507814	Kain M.	Kain
4	57197801981	Laurent B.	Laurent
5	57203867656	Leray F.	Leray
6	57203864793	Simon M.	Simon

Thus, if we retrieve a single author's information, we can gather other author IDs from this directly. If we have a specific paper, we can retrieve author IDs from that paper information as well.

Object retrieval

Along with metadata, abstract, and other information about a paper, additional objects may be accessed, such as figures, videos, or supplemental material. The `object_retrieval` function will return the objects associated with the identifier. The `process_object_retrieval` will convert the output into a tidy `data.frame`:

```
objects = object_retrieval("S1053811915002700", identifier = "pii", verbose = FALSE)
obj_df = process_object_retrieval(objects)
head(obj_df[, c("type", "url", "mime_type")])
```

	type
1	IMAGE-THUMBNAI
2	IMAGE-THUMBNAI
3	IMAGE-THUMBNAI
4	IMAGE-THUMBNAI
5	IMAGE-THUMBNAI
6	IMAGE-THUMBNAI

```

url
1 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr1.sml?httpAccept=image/gif
2 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr2.sml?httpAccept=image/gif
3 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr3.sml?httpAccept=image/gif
4 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr4.sml?httpAccept=image/gif
5 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr5.sml?httpAccept=image/gif
6 https://api.elsevier.com/content/object/eid/1-s2.0-S1053811915002700-gr6.sml?httpAccept=image/gif
  mime_type
1 image/gif
2 image/gif
3 image/gif
4 image/gif
5 image/gif
6 image/gif

```

where the URL provides the link to download the object. Here we will subset the output to high-resolution images, and pass the URL for the first to the `download_object` function. This will return the content of the image, as well as download the image to disk (with path outfile):

```

obj_df = obj_df[ grepl("image/jpeg", obj_df$mime_type),]
obj_df = obj_df[ obj_df$type %in% "IMAGE-HIGH-RES",]
object = download_object(obj_df$url[1])
object$outfile

[1] "/var/folders/1s/wrtqcpnx685_zk570bnx9_rr0000gr/T//Rtmp8s9r90/filedda42cd00281.jpg"

```

One can then view the output using the system viewer using `utils::browseURL`. The `download_objects` function is a convenience wrapper for `apply(url, download_object)` when passing in all the URLs from the output of `object_retrieval`.

Conclusion

The **rscopus** package provides an interface to the Scopus API through R. The package allows users to retrieve information about authors, individual articles, or institutions from the API. The information that can be obtained is limited by the scopes enabled by the API key through Elsevier and the associated organization. We have shown how to obtain an API key and specific use cases, which we believe covers the a number of the needs of the users.

More advanced usage requires additional information about the Scopus API and its querying specifications. The future direction of this package is to create helper functions to construct queries for more targeted views of the data. Additional APIs from Elsevier have functions for access, such as SciDirect with the `rscopus::sciencedirect_search` function, but the main focus of the package has been Scopus. Although the goal is to keep all functionality up to date with the number of APIs, the main focus will be consistency and stability with respect to Scopus. All the code and figures used to create this paper are located at <https://github.com/muschelli2/scopus> and the paper was created using the **rticles** package (1).

Funding

This research was supported by the Department of Biostatistics at Johns Hopkins Bloomberg School of Public Health.

Supplemental Material

Here is the simple parser `unique_title` to find the first 3 relevant words of the title after removing non-relevant words from a list in the **stopwords** package (3):

```

unique_title = function(x) {
  ss = sapply(strsplit(x, split = " "),
    function(x) {
      x = x[ !tolower(x) %in% stopwords::stopwords()]
      x = x[ !x %in% c("-", "?", "--", 1:100)]
      paste(x[1:3], collapse = " ")
    })
}

```

```

    })
    stopifnot(length(unique(ss)) == length(unique(x)))
    ss
  }

```

Bibliography

- [1] J. Allaire, Y. Xie, R Foundation, H. Wickham, Journal of Statistical Software, R. Vaidyanathan, Association for Computing Machinery, C. Boettiger, Elsevier, K. Broman, K. Mueller, B. Quast, R. Pruim, B. Marwick, C. Wickham, O. Keyes, M. Yu, D. Emaasit, T. Onkelinx, A. Gasparini, M.-A. Desautels, D. Leutnant, MDPI, O. Ögreden, D. Hance, and D. Nüst. *rticles: Article Formats for R Markdown*, 2018. URL <https://CRAN.R-project.org/package=rticles>. R package version 0.6. [p13]
- [2] M. Aria and C. Cuccurullo. *bibliometrix: An R-tool for comprehensive science mapping analysis*. *Journal of Informetrics*, 11(4):959–975, 2017. URL <https://doi.org/10.1016/j.joi.2017.08.007>. [p1]
- [3] K. Benoit, D. Muhr, and K. Watanabe. *stopwords: Multilingual Stopword Lists*, 2017. URL <https://CRAN.R-project.org/package=stopwords>. R package version 0.9.0. [p13]
- [4] L. Egghe. Theory and practise of the g-index. *Scientometrics*, 69(1):131–152, 2006. [p6]
- [5] L. Henry and H. Wickham. *purrr: Functional Programming Tools*, 2018. URL <https://CRAN.R-project.org/package=purrr>. R package version 0.2.5. [p8]
- [6] J. E. Hirsch. An index to quantify an individual’s scientific research output. *Proceedings of the National academy of Sciences*, 102(46):16569–16572, 2005. [p5]
- [7] J. Keirstead. *scholar: analyse citation data from Google Scholar*, 2016. URL <http://github.com/jkeirstead/scholar>. R package version 0.1.5. [p1]
- [8] J. Muschelli. *gcite: Google Citation Parser*, 2018. URL <https://CRAN.R-project.org/package=gcite>. R package version 0.9.2. [p1]
- [9] J. Ooms. The jsonlite package: A practical and consistent mapping between JSON data and R objects. *arXiv:1403.2805 [stat.CO]*, 2014. URL <https://arxiv.org/abs/1403.2805>. [p4]
- [10] H. Wickham. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York, 2016. ISBN 978-3-319-24277-4. URL <http://ggplot2.org>. [p5]
- [11] H. Wickham. *httr: Tools for Working with URLs and HTTP*, 2017. URL <https://CRAN.R-project.org/package=httr>. R package version 1.3.1. [p4]
- [12] H. Wickham and L. Henry. *tidyr: Easily Tidy Data with ‘spread()’ and ‘gather()’ Functions*, 2018. URL <https://CRAN.R-project.org/package=tidyr>. R package version 0.8.1. [p9]
- [13] H. Wickham, R. François, L. Henry, and K. Müller. *dplyr: A Grammar of Data Manipulation*, 2018. URL <https://CRAN.R-project.org/package=dplyr>. R package version 0.7.7. [p4]
- [14] H. Wickham et al. Tidy data. *Journal of Statistical Software*, 59(10):1–23, 2014. [p8]

John Muschelli

Department of Biostatistics, Johns Hopkins Bloomberg School of Public Health
615 N Wolfe St Baltimore, MD 21205
jmuschel@jhsph.edu